

Saarland University
Faculty of Mathematics and Computer Science
Department of Computer Science

Master's Thesis

Optimization-Based Precipitation Estimation from Commercial Microwave Links

submitted by

Iván Soto

Advisor

Prof. Dr.-Ing. Andreas Karrenbauer

Reviewers

Prof. Dr.-Ing. Andreas Karrenbauer (First Examiner)

Prof. Dr. Kurt Mehlhorn (Second Examiner)

August 2026

Eidesstattliche Erklärung Declaration of Original Authorship

Hiermit erkläre ich, dass ich die vorliegende Arbeit selbstständig und ohne unzulässige Hilfe verfasst habe und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe. Alle Stellen der Arbeit, die wörtlich oder sinngemäß aus Veröffentlichungen oder aus anderweitigen fremden Äußerungen entnommen wurden, sind als solche kenntlich gemacht. Ich habe KI-basierte Software, einschließlich OpenAI ChatGPT und OpenAI Codex, wie in der separaten Erklärung zur Nutzung generativer digitaler Assistenten beschrieben verwendet. Ich habe die eingereichte Arbeit geprüft und übernehme die volle Verantwortung für ihren Inhalt, einschließlich der Richtigkeit und Vollständigkeit ihrer Zitate und Literaturangaben. Ich erkläre mich damit einverstanden, dass meine Arbeit mithilfe von Plagiats- und KI-Erkennungssoftware überprüft werden darf. Mir ist bewusst, dass der Verstoß gegen diese Versicherung zum Nichtbestehen der Prüfung bis hin zum Verlust des Prüfungsanspruchs führen kann.

I hereby declare that this thesis is my own original work and was completed independently, without unauthorized assistance or unacknowledged sources. Any content derived from publications or other external sources, whether quoted verbatim or paraphrased, has been duly acknowledged and clearly marked as such. I used artificial intelligence-based software, including OpenAI ChatGPT and OpenAI Codex, as described in the separate Statement on the Usage of Generative Digital Assistants. I reviewed the submitted work and take full responsibility for its content, including the accuracy and completeness of its citations and references. I consent that my thesis may be checked using plagiarism and AI detection software. I acknowledge that any breach of this declaration may result in failing the examination and, in severe cases, the right to be examined may be revoked.

Saarbrücken, Datum/Date

Unterschrift/Signature

Einverständniserklärung (optional) Declaration of Consent (optional)

Ich bin damit einverstanden, dass meine bestandene Arbeit in beiden Versionen in die Bibliothek der Informatik aufgenommen und damit veröffentlicht wird.

I consent to having both the printed and digital versions of my thesis, which has been examined and awarded at least a passing grade, added to the Computer Science Department's library and thus made publicly accessible.

Saarbrücken, Datum/Date

Unterschrift/Signature

Abstract

Commercial microwave links provide an opportunistic source of rainfall information: rain attenuates the transmitted signal along each link path, and these attenuation measurements can be used to infer rainfall. Many existing CML rainfall-mapping approaches rely on link-level rain-rate estimates followed by spatial interpolation. In this thesis, we study optimization-based rainfall reconstruction from commercial microwave links as an alternative to purely interpolation-based map construction.

We construct a 100-patch benchmark from radar-derived European rainfall fields, simulate link attenuations using a physical attenuation model, and compare inverse-distance weighting (IDW), a link-aware interpolation variant (ILDW), and several optimization-based reconstructions. These include an ILDW-initialized nonlinear solver, a convex reformulation, a homotopy variant, and a ground-truth-initialized diagnostic solver. The evaluation compares pixelwise agreement with the radar-derived rainfall fields, distance-conditioned errors relative to link proximity, and agreement with the simulated link attenuations.

The results show that optimization-based reconstruction substantially improves upon IDW in patch-level RMSE and Pearson correlation. In particular, Convex Solver and Homotopy Solver reproduce the simulated link attenuations more accurately than both IDW and ILDW. At the same time, the case studies and attenuation-mismatch results show that minimizing the link-based objective is not always equivalent to recovering the radar-derived ground truth pixel by pixel. This highlights both the promise and the difficulty of formulating rainfall reconstruction from microwave-link measurements as an inverse optimization problem.

Acknowledgments

I am very grateful to Guy Even for wholeheartedly bringing me on board for this project, putting lots of effort into convincing me that this was something exciting to work on, spending so much time with me along the way to push it forward, and sharing so much of his wisdom in every meeting we had.

I also thank Andreas Karrenbauer for serving as the official advisor for this thesis as part of the requirements for submission of this work.

Contents

Abstract	i
Acknowledgments	iii
1 Introduction	1
2 Background and Related Work	3
3 Methodology	7
3.1 Data and Benchmark Construction	7
3.1.1 Geographic Distribution of the Benchmark Patches	8
3.1.2 Patch-Geometry Overview	9
3.2 Forward Model	9
3.3 Inverse Solvers	10
3.3.1 IDW	10
3.3.2 ILDW	11
3.3.3 Optimization-Based Reconstruction	12
4 Evaluation	17
4.1 Evaluation Protocol	17
4.1.1 Optimization Stopping Criteria	17
4.2 Patch-Level Metrics	18
4.2.1 Attenuation Mismatch	20
4.3 Distance-Conditioned Error Analysis	21
4.3.1 Rainy Pixel Relative Absolute Error by Link Distance	21
4.3.2 Non-Rainy Pixel Absolute Error by Link Distance	22
4.4 Wet/Dry Confusion Analysis	25
4.5 Reconstruction Runtime	26
5 Conclusion and Outlook	27
A Supplementary Analyses	29
A.1 Largest-Patch Link-Network Summary	29
A.2 Case Studies	30
Statement on the Usage of Generative Digital Assistants	33
References	35

Chapter 1

Introduction

Rainfall is one of the most important variables in meteorology, yet measuring it accurately over a geographic area remains difficult. One approach that has been studied extensively is to use commercial microwave links as an indirect source of rainfall observations (Messer et al., 2006; Goldshtein et al., 2009; Chwala and Kunstmann, 2019; Lian et al., 2022). These links, originally installed for telecommunications, are attractive because they are already widespread, operate continuously, and provide measurements that carry information about rainfall without requiring a dedicated rainfall-sensor network. Because rain attenuates the transmitted signal along the link path, attenuation measurements collected on such links can be used to infer rainfall.

Reconstructing rainfall fields from microwave-link data has already been studied in the literature, often through interpolation-based methods that first convert each link into a representative rain-rate value and then spread those values across a region (Lian et al., 2022). Simple interpolation rules such as inverse-distance weighting (IDW) are attractive because they are straightforward to implement and provide a practical baseline for rainfall mapping.

In this thesis, we study whether rainfall reconstruction can be improved by replacing purely interpolative map construction with optimization-based reconstruction. Rather than treating the rainfall map only as an interpolation problem, we formulate reconstruction as an optimization problem that balances agreement with the observed link attenuations against terms that encourage spatially coherent rainfall fields. We evaluate several solver variants derived from this formulation and compare them against IDW, the principal interpolation baseline used throughout this thesis.

To test this question, we construct a 100-patch benchmark from radar-derived rainfall fields. We then simulate microwave-link attenuations with a physical model relating rainfall to link attenuation and compare IDW, ILDW, and several solver variants derived from the optimization-based formulation.

With this setup in place, the remainder of the thesis reviews the background and prior work, presents the benchmark and methods, reports the empirical comparison, and closes with a discussion of the main conclusions and limitations.

For reproducibility, the code, experiment configuration files, and generated report artifacts used for this thesis are available in the project repository at <https://github.com/guyeven/Rainfall-Estimation>.

Chapter 2

Background and Related Work

In most settings, rainfall is measured primarily with rain gauges and weather radar. A rain gauge is a ground-based instrument that collects precipitation through an opening of known area and records the collected water automatically. It provides an accurate measurement at its installation site, but gauge networks are often too sparse to capture the spatial variability of rainfall across a region (McKee and Binns, 2016).

Weather radar takes a different approach by observing precipitation remotely. It transmits electromagnetic pulses and records the energy scattered back toward its antenna. The return time indicates the distance to the scattering particles, the antenna orientation indicates their direction, and the returned signal strength provides a measure called reflectivity. Radar therefore gives broad spatial coverage, but it does not measure rain rate directly; rain rate must be inferred from reflectivity using a conversion model. This inference is imperfect because the returned signal is also affected by the drop-size distribution, snow or hail, insects and birds, terrain reflections, partial obstruction of the radar beam by terrain or buildings, and other sources of error. Operational radar processing consequently requires filtering, classification, and calibration before reflectivity is converted into a rainfall product. Rainfall estimation can therefore remain challenging even where operational radar and gauge networks are already in place (Olsson et al., 2025). These limitations have motivated research into additional sources of rainfall information, including the attenuation measured along microwave communication links.

Microwave links as rainfall sensors. A microwave link is a fixed radio path between two antennas. Rain along the path absorbs and scatters part of the transmitted signal, reducing the power measured at the receiver. This reduction in received signal power is known as rain attenuation. Because rain attenuation degrades radio communication, much of the early propagation research focused on predicting this signal loss. At the same time, researchers recognized that the effect could be used for rainfall measurement. Atlas and Ulbrich (1977) described how microwave attenuation at suitable wavelengths could provide path- and area-averaged rainfall estimates. Olsen et al. (1978) subsequently used electromagnetic scattering calculations to derive relationships between rain rate and attenuation at different frequencies and compared the results with available measurements. The possibility of reconstructing a spatial rainfall field from several microwave paths was also explored before operational cellular-link data became available: Giuli et al. (1991) formulated a tomographic reconstruction method and evaluated it with simulated attenuation measurements from a purpose-designed network. Field experiments with dedicated microwave links later tested path-averaged rainfall retrieval in practice (Holt et al., 2000; Rahimi et al., 2003; Upton et al., 2005).

The shift to opportunistic sensing. The important change was then to reuse signal measurements from links already operating within cellular communication networks.

Messer et al. (2006) demonstrated the idea with operational links in Israel, while Leijnse et al. (2007) independently evaluated two operational links in the Netherlands against rain-gauge and radar observations. These studies provided complementary proofs of concept: their significance was not the number of links used, but the demonstration that existing cellular infrastructure could provide rainfall information without a dedicated meteorological link network. They also exposed practical uncertainties associated with communication-network data, including determination of the dry reference level, signal quantization, temporal sampling, and attenuation caused by water on the antennas.

From proof of concept to large-scale rainfall mapping. Research subsequently moved from estimating rainfall along individual operational links to reconstructing two-dimensional fields from irregular CML networks. Zinevich et al. (2008) formulated a tomographic method and evaluated it using attenuations simulated over the geometry and link characteristics of a real commercial network, while Goldshtein et al. (2009) developed another early network-based rainfall-field estimation method. A major increase in demonstrated operational scale came with country-wide rainfall maps for the Netherlands, produced from approximately 2,400 operational CMLs on selected rainy days (Overeem et al., 2013). These maps were constructed by assigning each path-averaged rainfall estimate to the midpoint of its link and interpolating the resulting point values with ordinary kriging. A later 2.5-year evaluation containing, on average, 3,383 CMLs showed that country-wide mapping could also be maintained over an extended period rather than only for selected events (Overeem et al., 2016).

The reconstruction problem. A CML observation belongs to an entire path, not to one point. Different rainfall distributions along and around that path can therefore produce the same total attenuation, so a network of links does not uniquely determine every pixel of a rainfall field. The interpretation also depends on link metadata. Each link has two endpoints, a carrier frequency, and a polarization; frequency and polarization determine the attenuation coefficients, while the endpoints determine which part of the rainfall field is sampled. Reconstructing a map consequently requires both a physical observation model and assumptions about spatial structure.

Reconstruction approaches. The historical studies above illustrate two broad strategies for constructing rainfall maps from CML observations. **Interpolation-based methods** first assign a path-averaged rain-rate estimate to each link and then interpolate those estimates across space. The country-wide studies used midpoint-based ordinary kriging, while **IDW-based methods** follow the same general two-stage structure with different interpolation weights (Overeem et al., 2013, 2016; Chwala and Kunstmann, 2019; Eshel et al., 2020). **Interpolation-based methods** first convert each link attenuation into a path-averaged rain-rate estimate and represent that estimate by one or more synthetic point observations. Once a link has been represented in this way, established methods for interpolating irregularly spaced point observations can be applied. The country-wide Netherlands studies used ordinary kriging with each link represented by its midpoint (Overeem et al., 2013, 2016), whereas Germany-wide mapping used midpoint-based inverse-distance weighting (IDW) (Graf et al., 2020). Eshel et al. (2021) applied both IDW and ordinary kriging to the same CML datasets and compared the resulting rainfall maps against radar reference fields. The Goldshtein–Messer–Zinevich (GMZ) algorithm provides a more link-aware interpolation variant: it places several virtual rain gauges along each link, adjusts their values iteratively so that their average remains consistent with the link-level estimate, and then applies IDW (Goldshtein et al., 2009; Eshel et al., 2021). These studies make IDW and ordinary kriging established interpolation references, with GMZ as a related link-aware

alternative. In the remainder of this thesis, IDW is the principal interpolation baseline; ordinary kriging and GMZ are discussed to place that choice within the broader reconstruction literature but are not evaluated in the benchmark. Inverse approaches instead retain the link paths and the physical attenuation relationship within the reconstruction. Zinevich et al. (2008) formulated nonlinear equations relating the rainfall field to the attenuation along each link. They solved these equations iteratively using Newton–Raphson and the simultaneous iterative reconstruction technique (SIRT), applying a smoothing operator during reconstruction. Their grid resolution adapted to the local link density, with smaller cells in areas containing more links. Bianchi et al. (2013) used an optimization method in a different setting, adjusting a radar rainfall field so that it better agreed with operational microwave-link and rain-gauge measurements.

In this thesis, we follow the inverse approach by estimating a rainfall field that agrees with the link attenuations while remaining spatially smooth. Our setup is closer to that of Zinevich et al. (2008): both approaches retain the complete link paths and seek a spatially smooth field that explains their attenuations. Their simulation study used an adaptive grid and estimated the spatial-correlation model controlling its smoothing from radar rainfall maps. We instead use a grid of uniform resolution and explicit objective terms for attenuation agreement and spatial smoothness, and compare several optimization variants with IDW across a 100-patch benchmark. For our deployable solvers, the radar-derived field is used to simulate the link observations and evaluate the reconstructions but is not supplied to the solver. The comparison with Bianchi et al. (2013) is less direct because their radar field was part of the reconstruction itself, whereas our optimization methods receive neither radar nor rain-gauge measurements. The methods are defined in detail in the methodology section.

The literature establishes CMLs as useful opportunistic rain sensors, but important methodological choices remain in how link measurements are turned into rainfall maps. In this thesis, we use a controlled benchmark for CML-based rainfall mapping to investigate whether optimization-based reconstruction can improve upon interpolation-based map construction.

Chapter 3

Methodology

3.1 Data and Benchmark Construction

Our benchmark is built patch by patch. A *patch* is a rectangular crop from a one-hour rainfall accumulation map in the OPERA-based EURADCLIM radar rainfall dataset¹ (Overeem et al., 2025), with accumulation expressed in millimetres. Each patch is selected to capture a detected rain event, with its width and height constrained to lie between 50 km and 95 km. The 100 patches are extracted from 84 distinct hourly EURADCLIM maps: 70 maps contribute one patch, 12 maps contribute two patches, and 2 maps contribute three patches. Because the patch detector is designed to capture rainy events rather than dry background, the benchmark is rain-dominated: across the 100 patches, the mean non-rainy proportion is about 10.50%, where a pixel is classified as non-rainy when its ground-truth rain rate is below 0.6 mm/h.

For each benchmark patch, we extract the corresponding rainfall crop from the EURADCLIM map. We then place this field on a finer 125 m grid by dividing each 2 km pixel into 16×16 subpixels, each assigned its parent pixel’s rainfall value. We smooth this field with a normalized Gaussian filter having a standard deviation of one refined pixel (125 m). The filter uses a 7×7 -pixel neighborhood, extending 375 m in each grid direction, and nearest-value padding at the patch boundaries. It redistributes rainfall among neighboring pixels, although the boundary padding means that the total rainfall is not necessarily preserved exactly. This pragmatic smoothing step softens the artificial block boundaries created by the subdivision. The resulting refined rainfall field serves as the benchmark ground truth for simulating link attenuations and evaluating the solvers.

The microwave-link geometry is not geographically co-located with the selected EURADCLIM patch locations. Instead, we use records from the 4TU-NL dataset (Overeem et al., 2024), which describes an operational CML network covering the Netherlands and provides the endpoint locations and operating frequencies of its links. We use these endpoint locations and frequencies as a realistic network template, place this network in the coordinate system of each selected patch, and retain only links whose endpoints fall inside the patch. Thus, the link geometries and frequencies are derived from actual Dutch CML infrastructure, whereas their placement relative to the EURADCLIM rainfall fields is synthetic. Across the 100 patches, this procedure retains between 939 and 1996 links

¹OPERA (Operational Programme for the Exchange of Weather Radar Information) is a programme of EUMETNET, a network of European national meteorological and hydrological services, that coordinates the exchange and compositing of weather-radar observations across Europe. EURADCLIM (European Climatological Gauge-Adjusted Radar Precipitation Dataset) was developed at the Royal Netherlands Meteorological Institute (KNMI). It is a processed historical precipitation dataset derived from OPERA radar composites that were further cleaned and merged with daily rain-gauge measurements from the European Climate Assessment & Dataset (ECA&D), a KNMI-coordinated database of quality-controlled weather-station observations contributed by data providers across Europe and the Mediterranean.

per patch (mean 1324). The attenuation of each retained link is then simulated from the radar-derived rainfall field rather than taken from the original operational CML measurements.

The 4TU-derived link file used in this benchmark does not specify polarization, so the patch-export pipeline supplies a default polarization when missing. In the present benchmark, that default is set to horizontal (H), and the forward model therefore uses the ITU coefficients corresponding to horizontal polarization.

To give a concrete sense of the benchmark, we next characterize the selected patches by their geographic distribution, dimensions, and retained 4TU-derived link counts.

3.1.1 Geographic Distribution of the Benchmark Patches

Figure 1 shows that the selected rainy-event patches are not concentrated in one country or local region; instead, they are spread across northern, western, central, southern, and southeastern Europe.



Figure 1: Geographic distribution of the 100 benchmark patches over the OPERA domain. Each red rectangle marks the footprint of one selected patch.

3.1.2 Patch-Geometry Overview




Figure 2: Patch width versus patch height for the 100 benchmark patches, with one dot per patch. Across patches, the mean patch area is 269,491 pixels with SD of 76,755 pixels (pixels are squares with side length 125 m).

Patch size also determines how many overlaid 4TU-derived links are retained, since only links with both endpoints inside the patch are used. Some retained links appear at more than one frequency, so Table 1 reports representative link counts and one- versus two-frequency counts for three patch sizes.

Patch size (km)	# links	1-multiplicity	2-multiplicity
50×50	939	99	420
62×60	1254	128	563
72×70	1507	157	675

Table 1: Representative link multiplicities for three patch dimensions. The multiplicity of a link is defined as the number of frequencies assigned to it. The second column reports the total number of link–frequency pairs. In the data set, link multiplicities were either 1 or 2, with multiplicity-one links accounting for approximately 10% of the reported link–frequency pairs.

3.2 Forward Model

In the benchmark, link attenuations are generated from the known benchmark ground-truth rainfall field using the forward model. Each link is represented as a line segment between its endpoints, and its attenuation is computed using ITU-R P.838-3, the standard relation for rain-induced specific attenuation in microwave links (International Telecom-

munication Union, 2005). In this model, the specific attenuation (dB/km) is

$$\gamma_R = \kappa(f, \text{pol}) R^{\alpha(f, \text{pol})},$$

where R is rain rate (mm/h), and (κ, α) depend on link frequency and polarization. For this calculation, each one-hour accumulation value is treated as the corresponding hourly mean rain rate. The path-integrated rain-induced attenuation along link ℓ , denoted by A_ℓ and hereafter referred to simply as the link attenuation, is computed as the line integral of the specific attenuation along the propagation path,

$$A_\ell \triangleq \int_\ell \gamma_R(R(s)) \, ds.$$

In the benchmark, and likewise for the methods we compare against, this integral is evaluated in discretized form on the refined 125 m grid: each link is traced through crossed grid cells, and attenuation is accumulated as a length-weighted sum of cellwise specific attenuations,

$$A_\ell \approx \sum_{c \in \mathcal{C}(\ell)} \kappa_\ell R_c^{\alpha_\ell} \Delta s_{\ell, c},$$

where $\mathcal{C}(\ell)$ is the set of crossed cells and $\Delta s_{\ell, c}$ is the path length of link ℓ inside cell c (in km). These computed attenuation values, together with each link's endpoint coordinates, frequency, and polarization, form the inputs to the inverse rainfall reconstruction.

3.3 Inverse Solvers

After the forward model defines how a rainfall field produces link attenuations, the reconstruction task runs in the opposite direction: given link-level observations, estimate the rainfall field that produced them. We compare IDW, ILDW, and four optimization-based reconstructions for this inverse problem: Solver(ILDW), Convex Solver, Homotopy Solver, and Solver(GT). Solver(GT) is initialized from the ground-truth rainfall field, so it is not a method that could be used on real unknown rainfall fields. Instead, we use it as an optimistic diagnostic reference for this objective and optimizer.

Using the same link-level observations and forward-model conventions described above, we now describe each solver family and its implementation details.

3.3.1 IDW

Inverse-distance weighting (IDW), originally introduced by Shepard for interpolating irregularly spaced point observations, is a spatial interpolation method (Shepard, 1968). In CML rainfall reconstruction, a common approach is to represent each link as a virtual rain gauge located at the link midpoint and interpolate its link-level rain-rate estimate across the spatial domain (Eshel et al., 2020). We use a truncated and discretized variant of this midpoint-IDW construction. The input consists of a set of links, each specified by its endpoints, rain-induced attenuation over the full link path A_ℓ , carrier frequency, and polarization.

IDW first converts each observed link attenuation A_ℓ into a single link-level rain-rate estimate defined as:

$$R_\ell \triangleq \left(\frac{A_\ell / |\ell|}{\kappa_\ell} \right)^{1/\alpha_\ell}.$$

where A_ℓ denotes the attenuation of link ℓ , $|\ell|$ its length, and $(\kappa_\ell, \alpha_\ell)$ the ITU-R P.838-3 coefficients corresponding to its frequency and polarization.

Each per-link rain-rate estimate is then treated as a virtual point observation at the link midpoint. For each pixel p with center c_p , observations within $r_{\max} = 15$ km are

assigned weights proportional to $\|m_\ell - c_p\|_2^{-2}$, corresponding to an IDW power parameter of 2; observations outside this radius receive zero weight. The weights of the retained observations are normalized to sum to one, and their weighted rain-rate estimates define $\hat{R}_p$. If no link midpoint lies within the search radius, $\hat{R}_p$ is set to zero. If one or more link midpoints lie inside pixel p , $\hat{R}_p$ is instead set to the mean of their rainfall estimates. This defines the truncated and discretized IDW baseline used throughout this thesis.

Algorithm 1: Discretized IDW. Each link is replaced with a virtual gauge located at the midpoint of the link. The rain rate at the gauge is set to the uniform rain rate that reproduces the input link attenuation under ITU-R P.838-3. IDW is then applied with respect to these virtual gauges. We set the exponent of the decay to 2.

Data: link set $\mathcal{L}$, a patch P , and a maximum influence distance $r_{\max}$

Result: rainfall field $\hat{R}$

```

1 foreach link  $\ell \in \mathcal{L}$  do                                     // compute rain rate per link
2    $(\ell.\kappa, \ell.\alpha) \leftarrow \text{ITU838}(\ell.\text{freq}, \ell.\text{polarization})$ 
3    $\ell.\text{rainrate} \leftarrow \left( \frac{\ell.\text{attenuation}/\ell.\text{length}}{\ell.\kappa} \right)^{1/\ell.\alpha}$ 
4 foreach pixel  $p \in P$  do                                     // compute rain  $\hat{R}_p$ 
5    $M_p \leftarrow \{\ell \in \mathcal{L} : \ell.\text{midpoint} \in p\}$ 
6   if  $M_p \neq \emptyset$  then
7      $\hat{R}_p \leftarrow \frac{1}{|M_p|} \sum_{\ell \in M_p} \ell.\text{rainrate}$            // pixel contains a link midpoint
8   else
9      $N_p \leftarrow \{\ell \in \mathcal{L} : \|\ell.\text{midpoint} - p.\text{center}\|_2 \leq r_{\max}\}$ 
10    if  $N_p = \emptyset$  then
11       $\hat{R}_p \leftarrow 0$                                        // pixel far from links
12    else
13       $\hat{R}_p \leftarrow \sum_{\ell \in N_p} \frac{\|\ell.\text{midpoint} - p.\text{center}\|_2^{-2}}{\sum_{\ell' \in N_p} \|\ell'.\text{midpoint} - p.\text{center}\|_2^{-2}} \cdot \ell.\text{rainrate}$ 

```

3.3.2 ILDW

ILDW is a link-aware modification of IDW used in this work. It starts from the same per-link rain-rate values obtained by inverting the discretized attenuation forward model for each link. The motivation for this variant is that a microwave-link attenuation measurement is associated with an entire path, not only with the midpoint of that path. We therefore considered whether interpolation could be improved by using the full link geometry when deciding which links should influence a pixel and how strongly they should do so.

More precisely, for a pixel center c_p , ILDW replaces the midpoint distance $\|m_\ell - c_p\|_2$ by the point-to-segment distance

$$\text{dist}(c_p, \ell),$$

where $\ell = [p_{\ell,1}, p_{\ell,2}]$ denotes the full link segment with endpoint coordinates $p_{\ell,1}$ and $p_{\ell,2}$.

If one or more links cross the pixel, the estimate is taken to be the mean of the corresponding link rain-rate values. Otherwise, the estimate is a weighted average over nearby links within radius $r_{\max}$ (set to 15 km), with weights proportional to the inverse squared segment distance. We considered this variant important because it preserves the same simple interpolation philosophy as IDW while using more of the link-path geometry.

At the algorithmic level, ILDW differs from Algorithm 1 only in how link geometry enters the interpolation step: midpoint membership and midpoint distance are replaced by pixel-segment intersection and distance to the full link segment.

Algorithm 2: Discretized ILDW. Link influence is determined using pixel-segment intersection and point-to-segment distance.

Data: link set $\mathcal{L}$, a patch P , and a maximum influence distance $r_{\max}$
Result: rainfall field $\hat{R}$

```

1 foreach link  $\ell \in \mathcal{L}$  do                                     // compute rain rate per link
2    $(\ell.\kappa, \ell.\alpha) \leftarrow \text{ITU838}(\ell.\text{freq}, \ell.\text{polarization})$ 
3    $\ell.\text{rainrate} \leftarrow \left( \frac{\ell.\text{attenuation}/\ell.\text{length}}{\ell.\kappa} \right)^{1/\ell.\alpha}$ 
4 foreach pixel  $p \in P$  do                                     // compute rain  $\hat{R}_p$ 
5    $X_p \leftarrow \{\ell \in \mathcal{L} : \ell \cap p \neq \emptyset\}$ 
6   if  $X_p \neq \emptyset$  then
7      $\hat{R}_p \leftarrow \frac{1}{|X_p|} \sum_{\ell \in X_p} \ell.\text{rainrate}$            // link crosses pixel
8   else
9      $N_p \leftarrow \{\ell \in \mathcal{L} : \text{dist}(p, \ell) \leq r_{\max}\}$ 
10    if  $N_p = \emptyset$  then
11       $\hat{R}_p \leftarrow 0$                                        // pixel far from links
12    else
13       $\hat{R}_p \leftarrow \sum_{\ell \in N_p} \frac{\text{dist}^{-2}(p.\text{center}, \ell)}{\sum_{\ell' \in N_p} \text{dist}^{-2}(p.\text{center}, \ell')} \cdot \ell.\text{rainrate}$ 

```

3.3.3 Optimization-Based Reconstruction

The optimization-based reconstructions take a different approach from interpolation. IDW and ILDW first reduce each link attenuation to a single link-level rain-rate estimate and then spread that information spatially by geometric weighting. In contrast, the optimization-based methods treat the rainfall value in every pixel as an unknown and choose the full rainfall field by minimizing an objective that combines agreement with the observed link attenuations with terms that encourage spatially coherent solutions.

Optimization Objective For each patch P , we identify P with its set of pixels, and denote by $\mathcal{L}(P)$ the set of links contained in P . Let $r \in \mathbb{R}_{\geq 0}^{|P|}$ denote the unknown rainfall field written as a vector over pixels. We solve the constrained optimization problem

$$\min_{r \in \mathbb{R}_{\geq 0}^{|P|}} J(r),$$

where the objective is a weighted sum of four terms:

$$J(r) \triangleq \alpha_{\text{atten}} J_{\text{atten}}(r) + \alpha_{1d} J_{1d}(r) + \alpha_{2d} J_{2d}(r) + \alpha_{\text{total}} J_{\text{total}}(r).$$

Here J_{atten} penalizes disagreement between observed and forward-model-predicted link attenuations, while J_{1d} , J_{2d} , and J_{total} are regularization terms that encourage the estimated rainfall field to remain spatially smooth and free of unnecessary oscillation, while weakly favoring smaller overall rainfall magnitudes. In the normalized solver studied here, we define the weights instance by instance using the corresponding objective values at the

ILDW solution:

$$\begin{aligned}\alpha_{\text{atten}} &\triangleq \frac{1}{J_{\text{atten}}(\text{ILDW})}, & \alpha_{1d} &\triangleq \frac{1}{J_{1d}(\text{ILDW})}, \\ \alpha_{2d} &\triangleq \frac{1}{2} \cdot \frac{1}{J_{2d}(\text{ILDW})}, & \alpha_{\text{total}} &\triangleq \frac{1}{|P|} \cdot \frac{1}{J_{\text{total}}(\text{ILDW})}.\end{aligned}$$

More concretely, J_{atten} is the attenuation-misfit term:

$$J_{\text{atten}}(r) \triangleq \frac{1}{|\mathcal{L}(P)|} \sum_{\ell \in \mathcal{L}(P)} \frac{(\hat{A}_\ell(r) - A_\ell^{\text{obs}})^2}{|\ell|},$$

where $\hat{A}_\ell(r)$ is the attenuation predicted by the forward model from the rainfall field r , A_ℓ^{obs} is the observed attenuation of link ℓ , and $|\ell|$ is the link length in km. Thus J_{atten} measures how far the forward-model-implied attenuations are from the observed attenuations, with a normalization by link length and averaging over the links in the patch.

In this benchmark, A_ℓ^{obs} is generated from the benchmark ground truth, but, except for the diagnostic Solver(GT), the reconstruction methods receive only these attenuation values and the corresponding link information. In an operational setting, A_ℓ^{obs} would instead be estimated from actual CML signal measurements after dry-baseline estimation and correction or filtering of non-rain effects, so the attenuation-misfit term would not require radar data as an input.

The term J_{1d} is a first-difference penalty:

$$J_{1d}(r) \triangleq \frac{1}{|P|} \sum_{(u,v) \in \mathcal{E}(P)} (r_u - r_v)^2,$$

where $\mathcal{E}(P)$ is the set of horizontally and vertically adjacent pixel pairs in patch P . This term sums squared differences between neighboring pixel values and therefore discourages abrupt pixel-to-pixel jumps in the estimated rainfall field.

The term J_{2d} is a second-difference penalty:

$$J_{2d}(r) \triangleq \frac{1}{|P|} \sum_{(p,q,t) \in \mathcal{T}(P)} ((r_q - r_p) - (r_t - r_q))^2,$$

where $\mathcal{T}(P)$ is the set of horizontal and vertical triples of collinear pixels in patch P . It penalizes changes in the first difference, so it discourages unnecessary curvature or oscillation along rows and columns.

Finally, J_{total} is an ℓ_2 shrinkage term on the rainfall field itself:

$$J_{\text{total}}(r) \triangleq \frac{1}{|P|} \sum_{p \in P} r_p^2.$$

This term biases the solution toward smaller overall rainfall values, except where larger values are needed to match the observed link attenuations.

These four terms are added together only after weighting, so the optimizer balances attenuation fit against smoothness, curvature control, and shrinkage.

These weights should be read as heuristic normalization choices rather than as parameters derived from an optimality principle. Using the ILDW objective values puts the four terms on comparable numerical scales at the same reference field, so that the optimizer is not dominated by one raw term simply because of units or magnitude. The additional factor 1/2 downweights the second-difference penalty relative to the first-difference penalty, while the factor $1/|P|$ makes the global shrinkage term a weak regularizer. We did not perform a systematic sensitivity analysis over nearby weight choices.

Solver Variants All optimization-based variants are solved iteratively with L-BFGS-B, a limited-memory quasi-Newton method for smooth constrained optimization (Byrd et al., 1995; Zhu et al., 1997). Starting from an initial rainfall field, the optimizer repeatedly updates the field to reduce the objective while enforcing the nonnegativity bounds on rainfall. In the present setting, the optimization variable is the rainfall value at every pixel of the patch. We write

$$\hat{r} \leftarrow \text{L-BFGS-B}(\mathcal{L}, P, r_{\text{init}})$$

to denote one invocation of L-BFGS-B on patch P with link set $\mathcal{L}$ and initial rainfall field r_{init} .

Solver(ILDW) The first variant applies L-BFGS-B directly to the original nonconvex objective above, using the ILDW rainfall field as the initial point:

$$\hat{r}_{\text{ILDW}} \leftarrow \text{L-BFGS-B}(\mathcal{L}, P, r_{\text{ILDW}}).$$

The objective is generally nonconvex because the ITU-R attenuation model uses a power law in rainfall. For a link ℓ , the predicted attenuation has the form

$$\hat{A}_\ell(r) = \sum_{c \in \mathcal{C}(\ell)} \kappa_\ell r_c^{\alpha_\ell} \Delta s_{\ell,c}.$$

When $\alpha_\ell = 1$, this predicted attenuation is linear in the pixel rain rates, so the squared attenuation residual is convex. For $\alpha_\ell \neq 1$, the attenuation residual is nonlinear and the resulting multivariable squared-residual term is not generally convex. As a simple scalar example, when $\alpha > 1$ and $A > 0$, $(r^\alpha - A)^2$ has negative curvature for sufficiently small positive r . Thus the solution found by L-BFGS-B can depend on the initial field.

We use ILDW as the initial point for this nonconvex solve because it is already constructed to be consistent with individual link attenuations in the simple case where links do not overlap. If links were pairwise disjoint at the pixel level, assigning each crossed pixel the rain-rate estimate of the link that crosses it would reproduce each observed attenuation exactly under the same forward model. Real patches contain overlapping links, so this exact consistency cannot generally hold for all links at once, but ILDW remains a natural initialization because it is organized around the same link paths whose attenuations the optimizer tries to fit.

Convex Solver The second variant constructs a convex surrogate inverse problem by replacing each original link with a virtual link. The virtual link keeps the same geometry and length, but uses a reference frequency f^0 for which the ITU-R attenuation exponent is one. In the present benchmark the links are treated as horizontally polarized, for which this reference frequency is $f^0 = 24.92$ GHz. For each original link, we first compute the uniform rain-rate equivalent

$$R_\ell \triangleq \left(\frac{A_\ell^{\text{obs}}}{|\ell| \kappa_\ell} \right)^{1/\alpha_\ell},$$

that is, the constant rain rate along the link that would reproduce its observed attenuation under the original link coefficients. The virtual observed attenuation is then

$$A_\ell^0 \triangleq |\ell| \kappa_0 R_\ell,$$

where κ_0 is the ITU-R coefficient at f^0 . We denote the resulting virtual link set by $\mathcal{L}^0$. Because the attenuation exponent is one for these virtual links, the predicted attenuation is linear in the rainfall field, so the attenuation-misfit term is a sum of squared linear

residuals. Together with the convex regularization terms, this gives a convex surrogate objective. We write

$$\hat{r}^0 \leftarrow \text{L-BFGS-B}(\mathcal{L}^0, P, r_{\text{init}})$$

for this solve. In the final experiments, r_{init} is the constant field 0.6 mm/h. This is the method reported as Convex Solver.

For Convex Solver, the attenuation weight is normalized within the virtual-link formulation:

$$\alpha_{\text{atten}}^{\text{Convex}} \triangleq \frac{1}{J_{\text{atten}}^{\text{virtual}}(r_{\text{ILDW}})},$$

where $J_{\text{atten}}^{\text{virtual}}$ is evaluated using the virtual link set $\mathcal{L}^0$ and virtual observed attenuations A_{ℓ}^0 . Thus, the ILDW rainfall field remains the reference field for normalization, but the denominator is its attenuation mismatch under the virtual-link model rather than $J_{\text{atten}}(r_{\text{ILDW}})$ for the original links.

Homotopy Solver The third variant uses homotopy continuation (Allgower and Georg, 2003) to move gradually from the convex surrogate problem back toward the original nonconvex formulation. For $\lambda \in [0, 1]$, define a virtual link set $\mathcal{L}^{\lambda}$ in which each link keeps the same geometry, length, and polarization as the original link, but has interpolated frequency

$$f_{\ell}^{\lambda} \triangleq (1 - \lambda)f^0 + \lambda f_{\ell}.$$

Let $\kappa_{\ell}(\lambda)$ and $\alpha_{\ell}(\lambda)$ be the ITU-R coefficients at frequency f_{ℓ}^{λ} . The corresponding virtual observed attenuation is

$$A_{\ell}^{\lambda} \triangleq |\ell| \kappa_{\ell}(\lambda) R_{\ell}^{\alpha_{\ell}(\lambda)}.$$

Starting from the convex problem at $\lambda = 0$, the solver advances through a sequence of increasing λ values, each time initializing from the previous solution:

$$\hat{r}^0 \leftarrow \text{L-BFGS-B}(\mathcal{L}^0, P, r_{\text{init}}), \quad \hat{r}^{\lambda} \leftarrow \text{L-BFGS-B}(\mathcal{L}^{\lambda}, P, \hat{r}^{\lambda-\delta}).$$

The final output is $\hat{r}^1$, which corresponds to the original link frequencies. The first stage uses the same constant 0.6 mm/h initial field as Convex Solver. At each subsequent stage, the previous stage's solution provides only the initial field. The attenuation term and its normalization correspond to the transformed link model at the current value of λ :

$$\alpha_{\text{atten}}^{\lambda} \triangleq \frac{1}{J_{\text{atten}}^{\lambda}(r_{\text{ILDW}})}.$$

Thus, the same ILDW rainfall field serves as the normalization reference at every stage, while $J_{\text{atten}}^{\lambda}$ is evaluated using that stage's link model and virtual observed attenuations. In the final experiments, this continuation-based method is reported as Homotopy Solver, and we use a continuation step size of $\delta = 0.1$.

Solver(GT) We also include a ground-truth-initialized diagnostic run:

$$\hat{r}_{\text{GT}} \leftarrow \text{L-BFGS-B}(\mathcal{L}, P, r_{\text{GT}}).$$

This run uses information unavailable in practice, so it is not a deployable reconstruction method. Instead, it provides an optimistic empirical reference for interpreting the objective landscape and for seeing what the objective preserves, or moves away from, when initialized at the target rainfall field.

Chapter 4

Evaluation

4.1 Evaluation Protocol

We evaluate each solver in four ways: overall agreement with the ground-truth rainfall map, how errors change as distance from links increases, wet/dry classification behavior, and agreement with observed link attenuations under the forward model. We first report patch-level summaries across all 100 benchmark patches, and then provide qualitative case studies in the appendix. We additionally report end-to-end runtime for IDW, ILDW, and Convex Solver to compare the time required by each method to reconstruct a patch. All optimization-based results reported below use the objective terms and normalization rules described above. The single-stage variants have a maximum budget of 2000 L-BFGS-B iterations per patch, while the Homotopy Solver applies this limit separately at each continuation stage.

4.1.1 Optimization Stopping Criteria

The 2000-iteration budget was used as a cap rather than as a fixed number of iterations. Every L-BFGS-B invocation, including each stage of the Homotopy Solver, used the same stopping tolerances, with relative objective-reduction tolerance $\text{ftol} = 10^{-6}$ and projected-gradient tolerance $\text{gtol} = 10^{-6}$. Let J_k denote the objective value at iteration k . The ftol criterion stops the solve when the relative objective decrease between two successive iterations is small,

$$\frac{J_k - J_{k+1}}{\max\{|J_k|, |J_{k+1}|, 1\}} \leq \text{ftol}.$$

The gtol criterion stops the solve when the projected gradient is small,

$$\|g_{k+1}^{\text{proj}}\|_{\infty} \leq \text{gtol},$$

where g_{k+1}^{proj} is the gradient projected according to the bound constraints used by L-BFGS-B. The ftol criterion establishes that the relative objective reduction has become small, whereas satisfying gtol provides the stronger indication that the bound-constrained optimizer is close to a stationary point. No Solver(ILDW), Convex Solver, or Solver(GT) reconstruction reached its 2000-iteration cap, and no Homotopy Solver stage reached the corresponding per-stage cap. Thus, none of the reported reconstructions terminated because it exhausted the configured iteration budget. Because nearly all runs stopped through ftol , these stop reasons indicate negligible objective change between successive iterations, but not necessarily a small projected gradient.

Solver	Mean iter.	SD iter.	Min	Max	Final stop reason
Solver(ILDW)	557.4	289.1	3	1198	100 ftol
Convex Solver	936.7	233.8	387	1557	99 ftol, 1 gtol
Homotopy Solver	1072.4	243.3	429	1712	100 ftol
Solver(GT)	594.4	345.9	3	1787	100 ftol

Table 2: Iteration counts for the optimization-based reconstruction runs across the 100 evaluated patches. For the single-stage variants, the count is the number of iterations in that patch’s L-BFGS-B invocation. For Homotopy Solver, it is the total number of inner L-BFGS-B iterations summed across all continuation stages, while the displayed stop reason refers to the final stage. The standard deviation is computed across patches. No single-stage reconstruction or Homotopy stage reached its applicable 2000-iteration cap.

4.2 Patch-Level Metrics

We now turn to the empirical comparison on the hundred-patch benchmark. We begin with whole-patch map metrics and attenuation consistency, then examine how pixel-level errors vary with link proximity, followed by wet/dry classification and runtime.

For a patch P , let $R_P(p)$ denote the ground-truth rainfall at pixel p , and let $\hat{R}_{P,s}(p)$ denote the rainfall predicted at pixel p by solver s . We compute three scalar metrics over all pixels in P : the root mean squared error

$$\text{RMSE}_{P,s} \triangleq \sqrt{\frac{1}{|P|} \sum_{p \in P} (\hat{R}_{P,s}(p) - R_P(p))^2},$$

the signed bias

$$\text{Bias}_{P,s} \triangleq \frac{1}{|P|} \sum_{p \in P} (\hat{R}_{P,s}(p) - R_P(p)),$$

and the Pearson correlation coefficient

$$r_{P,s} \triangleq \frac{\sum_{p \in P} (\hat{R}_{P,s}(p) - \bar{\hat{R}}_{P,s})(R_P(p) - \bar{R}_P)}{\sqrt{\sum_{p \in P} (\hat{R}_{P,s}(p) - \bar{\hat{R}}_{P,s})^2} \sqrt{\sum_{p \in P} (R_P(p) - \bar{R}_P)^2}},$$

where

$$\bar{\hat{R}}_{P,s} \triangleq \frac{1}{|P|} \sum_{p \in P} \hat{R}_{P,s}(p), \quad \bar{R}_P \triangleq \frac{1}{|P|} \sum_{p \in P} R_P(p).$$

RMSE measures overall magnitude error. Bias measures average signed error (net over- or underestimation). Correlation measures how well the spatial structure of the field is reproduced independently of scale and offset. The Pearson correlation coefficient satisfies $r_{P,s} \in [-1, 1]$. A value of 1 indicates perfect positive linear agreement between predicted and ground-truth rainfall maps. A value of 0 indicates no linear relationship. A value of -1 indicates perfect negative linear agreement (an exact inverse pattern).

Solver	RMSE (mm/h)	Signed bias (mm/h)	Correlation
IDW	1.307 ± 1.234	0.034 ± 0.296	0.866 ± 0.059
ILDW	1.239 ± 1.182	0.047 ± 0.261	0.878 ± 0.055
Solver(ILDW)	1.059 ± 1.192	0.040 ± 0.235	0.918 ± 0.055
Convex Solver	0.845 ± 0.803	0.023 ± 0.075	0.940 ± 0.034
Homotopy Solver	0.845 ± 0.803	0.023 ± 0.075	0.940 ± 0.034
Solver(GT)	0.503 ± 0.369	0.009 ± 0.030	0.961 ± 0.032

Table 3: Descriptive patch-level summary of three global map metrics across the 100 evaluated patches. Each entry reports mean $\pm$ standard deviation across patches. Here RMSE and signed bias are computed in the same rainfall units as R_P and $\hat{R}_{P,s}$, while Correlation is the Pearson correlation coefficient between the predicted and ground-truth rainfall fields within each patch.

Table 3 gives a descriptive overview of the scale and variability of the patch-level metrics for each solver. Homotopy Solver gives nearly identical descriptive metrics to Convex Solver, so in these experiments its added continuation machinery does not produce an additional empirical gain in the map-level summaries. Solver(GT) has the lowest mean RMSE, lowest mean signed bias, and highest mean correlation, as expected for a run initialized from the ground truth. Comparing IDW and ILDW also suggests that keeping more of the link-path geometry can be useful even within a simple interpolation framework: ILDW has lower mean RMSE and higher mean correlation than midpoint IDW, although its mean signed bias is not lower than IDW’s.

Because all solvers are evaluated on the same 100 patches, direct solver comparisons can use paired patchwise differences. We treat these differences as approximately independent because the patches are distinct spatial crops and are often widely separated, even when they come from the same hourly EURADCLIM map. Among the 18 pairs of patches extracted from the same hourly EURADCLIM map, the median distance between patch centers is approximately 544 km, and only two pairs overlap spatially. The independence assumption may not hold exactly, however. Considering patches extracted from different hourly maps, 16 pairs both share at least one native EURADCLIM grid cell and occur within six hours of each other; six of these pairs are separated by exactly one hour. These pairs could represent successive observations of the same evolving rainfall system, and their solver differences could therefore be correlated.

We restrict the paired comparison to Convex Solver and IDW, the principal optimized method and interpolation baseline, respectively. Convex Solver has lower RMSE and higher Pearson correlation on all 100 patches. For absolute bias, Convex Solver is closer to zero on 78 patches and IDW is closer on 22 patches.

Unlike the RMSE and correlation comparisons, the absolute-bias comparison does not favor Convex Solver for every patch. We therefore use a one-sided paired t -test, implemented as a one-sample t -test on the patch-level paired differences, to assess whether Convex Solver has lower absolute bias than IDW on average rather than interpreting the win counts alone. Because this test assumes that the patch-level paired differences are independent, its p -value should be interpreted cautiously in light of the possible spatiotemporal dependence noted above.

For absolute bias, we define

$$d_P = |\text{Bias}_{P,\text{Convex}}| - |\text{Bias}_{P,\text{IDW}}|,$$

where a negative value means that Convex Solver is closer to zero bias on patch P . We test $H_0 : \mathbb{E}[d_P] \geq 0$ against $H_1 : \mathbb{E}[d_P] < 0$ using the 100 patch-level differences. The test rejects the null hypothesis at $\alpha = 0.01$ ($p = 4.667 \times 10^{-7}$).

Metric	Better direction	Convex wins	IDW wins	Patch-level evidence
RMSE	lower	100	0	all patches favor Convex Solver
Absolute Bias	lower	78	22	rejects null: $p = 4.667 \times 10^{-7}$
Pearson correlation	higher	100	0	all patches favor Convex Solver

Table 4: Paired patchwise comparison between Convex Solver and IDW across the 100 evaluated patches.

4.2.1 Attenuation Mismatch

The previous metrics compare each reconstructed rainfall map directly to the ground-truth radar field. We also evaluate the reconstructions from the point of view of the CML observations themselves: if a predicted rainfall field is passed through the forward attenuation model, how closely does it reproduce the observed link attenuations? This matters because a useful reconstruction should not only resemble the radar field pixelwise, but should also remain consistent with the link observations that drive the inverse problem.

We report two complementary attenuation-mismatch summaries in the same table. The mean absolute attenuation error per kilometer gives a directly interpretable link-level discrepancy in dB/km, while J_{atten} gives a squared attenuation-mismatch measure. For every solver, both measures are recomputed at each link’s benchmark frequency and compared with the corresponding observed attenuation. Thus, J_{atten} is a shared post-hoc evaluation measure rather than necessarily the attenuation term optimized internally by a given solver. In particular, Convex Solver optimizes $J_{\text{atten}}^{\text{virtual}}$, while Homotopy Solver successively optimizes $J_{\text{atten}}^{\lambda}$ for multiple values of λ ; the final reconstructions from both methods are evaluated here using J_{atten} .

Solver	Abs. err. mean (dB/km)	Abs. err. SD (dB/km)	J_{atten} mean (dB ² /km)	J_{atten} SD (dB ² /km)
IDW	0.05034	0.03396	0.05284	0.13575
ILDW	0.01356	0.00943	0.004204	0.01390
Solver(ILDW)	0.00207	0.00373	0.0002951	0.001466
Convex Solver	0.00107	0.00099	0.00003354	0.0001431
Homotopy Solver	0.00075	0.00068	0.00001333	0.00005245
Solver(GT)	0.00129	0.00166	0.00004951	0.0001492

Table 5: Patch-level attenuation-mismatch summaries across the 100 evaluated patches.

For a given patch and solver s , the absolute-error columns summarize

$\frac{1}{N_{\text{links}}} \sum_{\ell \in \text{links}} \frac{|\hat{A}_{\ell,s} - A_{\ell}^{\text{obs}}|}{|\ell|}$, while $J_{\text{atten}} = \frac{1}{N_{\text{links}}} \sum_{\ell \in \text{links}} \frac{(\hat{A}_{\ell,s} - A_{\ell}^{\text{obs}})^2}{|\ell|}$. Here, $\hat{A}_{\ell,s}$ is recomputed from the rainfall field estimated by solver s at link ℓ ’s benchmark frequency, A_{ℓ}^{obs} is the observed attenuation, and $|\ell|$ is the link length in kilometers. Lower values indicate better agreement with the observed link attenuations.

Together, these attenuation-based summaries show that the optimized solvers match the observed link attenuations much more closely than IDW and ILDW, with Convex Solver and Homotopy Solver giving the smallest mismatches. They also add a useful qualification to the map-level results above: pixelwise agreement with the radar-derived ground truth and consistency with the simulated CML observations are two measures that capture different aspects of reconstruction quality. Solver(GT) remains closest to the radar field, but Convex Solver and Homotopy Solver better satisfy the link observations under the forward model. Thus, the field closest to the radar ground truth is not necessarily the field that minimizes attenuation mismatch. In addition, although Homotopy Solver does

not improve the RMSE, bias, or correlation summaries relative to Convex Solver, it does provide a small gain in link consistency.

4.3 Distance-Conditioned Error Analysis

We next ask how prediction quality varies with link proximity. To answer this question, we group pixels into distance bins according to their distance to the third-closest link and summarize patch-level error statistics within each bin. We use distance to the third-closest link as a proxy for local link density: small values indicate several nearby links, whereas large values indicate sparse coverage. Unlike nearest-link distance, this measure does not classify a location as densely covered solely because one link is nearby. The choice of three is a practical analysis choice rather than a value dictated by the physical model.

We examine rainy pixels using relative absolute error and non-rainy pixels using absolute error. In this subsection we focus exclusively on distance-conditioned pixelwise error summaries.

4.3.1 Rainy Pixel Relative Absolute Error by Link Distance

We define a pixel p to be rainy if $R_P(p) \geq 0.6$ mm/h, and we write

$$P_{\text{rain}} = \{p \in P : R_P(p) \geq 0.6\}.$$

We use 0.6 mm/h as a pragmatic analysis threshold to exclude trace and very weak rainfall, for which relative errors are highly sensitive to small absolute deviations and operational CML detection becomes less reliable (Polz et al., 2020). The threshold does not represent a universal physical boundary between rain and no rain, so the wet/dry summaries below should still be interpreted as threshold-dependent.

For a rainy pixel p , the relative absolute error with respect to solver s is

$$\text{RAE}_{P,s}(p) \triangleq \frac{|R_P(p) - \hat{R}_{P,s}(p)|}{R_P(p)}.$$

To study how prediction quality changes with link proximity, we group pixels by distance to the third-closest link. For a pixel $p \in P$, let $c(p) \in \mathbb{R}^2$ denote the center of pixel p . For a link $\ell \in \mathcal{L}(P)$, viewed geometrically as a line segment, let

$$\text{dist}(p, \ell) \triangleq \text{dist}(c(p), \ell)$$

denote the Euclidean distance from the pixel center to that link segment. If the links in patch P are written as $\ell_1, \dots, \ell_{|\mathcal{L}(P)|}$, we define

$$d_3(p) \triangleq \text{the third-order statistic of } \{\text{dist}(p, \ell_j)\}_{j=1}^{|\mathcal{L}(P)|},$$

that is, the distance from the center of pixel p to its third-closest link segment.²

For each index i , let b_{i-1} and b_i denote the lower and upper distance cutoffs of the i th bin, and write $B_i = (b_{i-1}, b_i]$. A pixel p belongs to bin B_i exactly when $d_3(p) \in B_i$.

For each patch P , solver s , and distance bin B_i , we then collect the rainy-pixel RAE values in that bin and summarize them by the patch-level median

$$m_{s,i}^{\text{rain}}(P) \triangleq \text{median}\{\text{RAE}_{P,s}(p) : p \in P_{\text{rain}}, d_3(p) \in B_i\}.$$

²For computational efficiency, $d_3(p)$ is approximated using a sampled candidate search, followed by exact point-to-segment distance calculations for the candidate links.

If a patch contains no rainy pixels in a particular distance bin, this median is undefined and that patch does not contribute to the summary for that bin. Consequently, the number of contributing patches may differ between bins.

In Figure 3, these patch-level medians are compared across the full benchmark collection $\mathcal{P}$ of 100 patches. There, the displayed distance bins are

$$\begin{aligned} &[0, 125] \text{ m}, \quad (125, 375] \text{ m}, \quad (375, 750] \text{ m}, \\ &\quad (750, 1500] \text{ m}, \quad (1500, 3125] \text{ m}, \\ &\quad (3125, 6000] \text{ m}, \quad (6000, \infty) \text{ m}. \end{aligned}$$

In the box-and-whisker plot, the two sparsely populated tail bins $(6000, 9000] \text{ m}$ and $(9000, \infty)$ are merged and displayed as a single $(6000, \infty)$ m bin. The count table below still reports these two tail bins separately.

For fixed solver s and distance bin B_i , let

$$\mu_{s,i}^{\text{rain}} \triangleq \{m_{s,i}^{\text{rain}}(P) : P \in \mathcal{P}, m_{s,i}^{\text{rain}}(P) \text{ is defined}\}$$

denote the sequence of patch-level rainy-pixel median RAE values across patches. Each box-and-whisker then summarizes this sequence: the lower whisker tip is the minimum of $\mu_{s,i}^{\text{rain}}$, the lower edge of the box is its 25th percentile, the horizontal line segment with the filled circle at its center marks the median, the upper edge of the box is its 75th percentile, and the upper whisker tip is its maximum.




Figure 3: Rainy-pixel distance-profile box-and-whisker summary for distance to the third closest link. Each patch containing at least one rainy pixel in a given distance bin contributes one median rainy-pixel relative-absolute-error (RAE) value for that bin, and these patch-level medians are then summarized across patches with a box-and-whisker plot. Rainy pixels are those whose ground-truth rainfall is at least the configured wet threshold, set here to 0.6 mm/h. The corresponding rainy-pixel counts are reported separately in Table 6.

4.3.2 Non-Rainy Pixel Absolute Error by Link Distance

For the dry-pixel analysis, we focus on pixels whose ground-truth rainfall falls below the wet threshold. We therefore define the non-rainy pixel set by

$$P_{\text{nonrain}} = \{p \in P : R_P(p) < 0.6\}.$$

Distance to third closest link (m)	Rainy mean count (% of rainy pixels)	SD
[0, 125]	4267.0 (1.79%)	960.4
(125, 375]	16737.7 (7.01%)	3803.2
(375, 750]	26809.9 (11.22%)	6707.0
(750, 1500]	51988.5 (21.76%)	14 468.0
(1500, 3125]	88839.4 (37.18%)	27 531.1
(3125, 6000]	45929.4 (19.22%)	15 679.4
(6000, 9000]	3835.8 (1.61%)	2725.7
(9000, ∞)	514.4 (0.22%)	966.9
Total	238922.1 (100%)	

Table 6: Rainy-pixel counts per distance bin for the same distance-to-third-closest-link partition used in Figure 3. For each patch, pixels are classified by the value of $d_3(p)$, and the table reports the mean and standard deviation of the rainy-pixel count in each bin across patches.

For these pixels, relative normalization by $R_P(p)$ is not appropriate because the denominator may be zero or very small. We therefore measure prediction quality by absolute error rather than relative absolute error:

$$\text{AE}_{P,s}(p) = |R_P(p) - \hat{R}_{P,s}(p)|.$$

Using the same definition of the third-closest-link distance $d_3(p)$ introduced above, and the same distance bins B_i as in the rainy-pixel analysis, we compute for each patch P , solver s , and distance bin B_i

$$m_{s,i}^{\text{nonrain}}(P) \triangleq \text{median}\{\text{AE}_{P,s}(p) : p \in P_{\text{nonrain}}, d_3(p) \in B_i\},$$

and, for fixed s and i , the corresponding box-and-whisker plot summarizes the sequence

$$\mu_{s,i}^{\text{nonrain}} \triangleq \{m_{s,i}^{\text{nonrain}}(P) : P \in \mathcal{P}, m_{s,i}^{\text{nonrain}}(P) \text{ is defined}\}$$

of patch-level non-rainy-pixel median AE values across patches. As in the rainy-pixel analysis, patches without eligible pixels in a bin are omitted from that bin’s summary. Thus the rainy and non-rainy figures use the same distance-bin construction, but they differ both in the pixel set being summarized and in the error metric being aggregated.




Figure 4: Non-rainy-pixel distance-profile box-and-whisker summary for distance to the third closest link. Each patch containing at least one non-rainy pixel in a given distance bin contributes one median non-rainy-pixel absolute error value for that bin, and these patch-level medians are then summarized across patches with a box-and-whisker plot. Non-rainy pixels are those whose ground-truth rainfall is below the configured wet threshold, set here to 0.6 mm/h. The corresponding non-rainy-pixel counts are reported separately in Table 7.

Distance to third closest link (m)	Non-rainy mean count (% of non-rainy pixels)	SD
[0, 125]	563.6 (1.84%)	777.6
(125, 375]	2261.1 (7.40%)	3093.5
(375, 750]	3389.1 (11.09%)	4709.5
(750, 1500]	6446.5 (21.09%)	9257.7
(1500, 3125]	10776.2 (35.25%)	15 553.7
(3125, 6000]	6197.2 (20.27%)	8990.0
(6000, 9000]	849.9 (2.78%)	1289.7
(9000, ∞)	85.5 (0.28%)	250.9
Total	30569.2 (100%)	

Table 7: Non-rainy-pixel counts per distance bin for the same distance-to-third-closest-link partition used in Figure 4. For each patch, pixels are classified by the value of $d_3(p)$, and the table reports the mean and standard deviation of the non-rainy-pixel count in each bin across patches.

To make the distance-conditioned summaries easier to interpret geometrically, we show the distance-bin assignment for the largest evaluated patch. This illustrates how the third-closest-link distance partitions a real patch into near-link and low-coverage regions.




Figure 5: Distance-bin map for the largest evaluated patch in the data set. Each pixel is colored according to the bin induced by its distance to the third closest link, and the complete link geometry is overlaid in white.

4.4 Wet/Dry Confusion Analysis

Beyond pixel-level error magnitudes, we also summarize whether each method preserves the basic wet/dry structure of the rainfall field. Using the same 0.6 mm/h threshold as above, each pixel is classified as wet or dry in both the ground truth and the reconstruction. This gives patch-level true-positive, false-negative, false-positive, and true-negative fractions, summarized in Table 8.

Solver	TP/ $ P $	FN/ $ P $	FP/ $ P $	TN/ $ P $
IDW	0.8926 ± 0.0126	0.0024 ± 0.0004	0.0396 ± 0.0039	0.0654 ± 0.0096
ILDW	0.8921 ± 0.0127	0.0029 ± 0.0004	0.0370 ± 0.0037	0.0680 ± 0.0097
Solver(ILDW)	0.8919 ± 0.0127	0.0031 ± 0.0004	0.0324 ± 0.0036	0.0726 ± 0.0098
Convex Solver	0.8912 ± 0.0127	0.0038 ± 0.0005	0.0228 ± 0.0022	0.0823 ± 0.0109
Homotopy Solver	0.8912 ± 0.0127	0.0038 ± 0.0005	0.0228 ± 0.0022	0.0822 ± 0.0109
Solver(GT)	0.8929 ± 0.0125	0.0021 ± 0.0003	0.0114 ± 0.0013	0.0936 ± 0.0123

Table 8: Patch-averaged wet/dry confusion matrix entries at a threshold of 0.6 mm/h, where a pixel is called wet if its rainfall is at least 0.6 mm/h and dry otherwise. For each patch, pixels are first assigned to the four confusion categories defined by the ground-truth and predicted wet/dry labels. The count in each category is then divided by the total number of pixels in that patch, producing a patch-level pixel fraction. The table entries report the mean of those patch-level fractions across the 100 evaluated patches, together with the standard error of the mean (SEM).

False negatives warrant particular attention because they correspond to missed rain, while false positives create spurious wet areas. Wet/dry rainfall classification generally

involves a trade-off between these errors, and their relative importance depends on the intended application (Polz et al., 2020). Because these rain-event patches contain approximately 89.5% wet pixels on average, normalizing all four entries by the total patch size makes the false-positive fractions shown in Table 8 appear relatively small. Conditional on ground-truth dry pixels, the mean of the patch-level conditional false-positive rates is approximately 59.8% for IDW and 40.0% for Convex Solver. Conditional on ground-truth wet pixels, the mean of the corresponding patch-level conditional false-negative rates remains small, at approximately 0.32% and 0.50%, respectively, meaning that both methods identify at least 99.5% of ground-truth wet pixels. Thus, Convex Solver remains close to IDW in wet-pixel detection, with a slightly higher false-negative rate, while substantially reducing false wet-area detection.

4.5 Reconstruction Runtime

Having compared reconstruction quality across the preceding analyses, we next assess the computational feasibility of Convex Solver for near-real-time applications. We timed each of the 100 benchmark patches using Convex Solver, with IDW and ILDW included for comparison. This focused experiment was not intended as an exhaustive runtime comparison across all solver variants; Solver(ILDW), Homotopy Solver, and the diagnostic Solver(GT) were therefore not timed. The patches and methods were processed sequentially on an Apple M2 MacBook Air with an 8-core CPU and 8 GB of memory. Each measurement is an end-to-end wall-clock time that includes input loading, initialization, reconstruction, and writing the reconstructed rainfall map and associated solver outputs. The resulting runtime summaries are reported in Table 9.

Solver	Mean $\pm$ SD (s)	Median (s)	Min-max (s)
IDW	7.53 ± 2.43	7.10	4.01–14.18
ILDW	18.69 ± 6.20	18.25	8.67–38.19
Convex Solver	76.03 ± 27.08	71.31	32.76–153.12

Table 9: End-to-end wall-clock reconstruction times across the 100 benchmark patches. The methods and patches were run sequentially on an Apple M2 MacBook Air with an 8-core CPU and 8 GB of memory.

Summary of findings. Across the benchmark, the optimized solvers have lower patch-level RMSE and higher Pearson correlation than IDW. Convex Solver and Homotopy Solver also give stronger rainfall-map agreement than Solver(ILDW) while maintaining strong attenuation consistency. Homotopy Solver offers no map-level improvement over Convex Solver, only a small gain in attenuation consistency. Solver(GT) has the strongest map-level summaries, whereas Convex Solver and Homotopy Solver achieve lower attenuation mismatch, confirming that the map-level and link-based measures capture different aspects of reconstruction quality.

Errors generally increase as nearby-link coverage becomes weaker, but the optimized solvers retain lower median errors than IDW in the low-coverage bins. Convex Solver reduces false wet-area detection relative to IDW at the cost of a small increase in missed-rain pixels. Its 76.03-second mean runtime suggests potential for near-real-time use at the benchmark scale.

Chapter 5

Conclusion and Outlook

This thesis compared interpolation and optimization-based reconstruction strategies on a 100-patch benchmark constructed from the OPERA-based EURADCLIM dataset and simulated microwave-link attenuations. In the patch-level attenuation-mismatch summaries, every optimization-based solver had lower mean mismatch than IDW on both reported measures, computed at the benchmark link frequencies. Convex Solver and Homotopy Solver also had lower map-level RMSE and higher Pearson correlation summaries than both IDW and the ILDW-initialized nonlinear solver. In the paired comparison with IDW, Convex Solver also had lower absolute patch-level bias on average. Solver(GT) had the strongest reported map-level summaries, but it is initialized from the ground truth and therefore is not a deployable method. At the same time, the comparison between Solver(GT) and the convex solvers shows that the rainfall field closest to the radar-derived ground truth is not always the field that best matches the link observations under the chosen forward model. ILDW remained a useful link-aware IDW modification: it preserves the simplicity of IDW-style interpolation while incorporating the full link path into the interpolation weights rather than representing each CML observation only at its midpoint.

An important next step is to repeat the comparison of reconstruction methods using measurements from an operational CML network. In the current benchmark, the observed attenuations are generated from the radar-derived ground-truth rainfall fields using the benchmark forward model, producing noise-free and model-consistent observations. In a real-time application, the attenuation-mismatch term would instead be driven by rain-induced attenuations estimated from current CML signal measurements, while independent radar or rain-gauge observations would be withheld from the reconstruction and used only afterward as reference data for evaluating the resulting rainfall estimates. Such an experiment is needed to determine whether the improvements over IDW persist in the presence of measurement noise, changing link availability, wet-antenna attenuation, baseline drift, and mismatch between the forward model and real signal behavior.

Two further aspects of the controlled benchmark warrant validation. First, all patches use the same 4TU-derived network template, so the results do not establish generalization across other network geometries or densities. Second, the benchmark target is obtained by upsampling and Gaussian smoothing of the coarse EURADCLIM rainfall fields, while the optimization objective explicitly favors spatial smoothness. This alignment may benefit the regularized solvers. Future work should therefore repeat the comparison across varied operational network layouts and with rainfall references that do not rely on the same subdivision-and-smoothing procedure.

Operational applicability also depends on computational latency. As reported in Table 9, Convex Solver required an average of 76.03 seconds per patch on the benchmark hardware (median 71.31 seconds; range 32.76–153.12 seconds). Whether this runtime is compatible with near-real-time reconstruction depends on the target update interval and on end-to-end operational latency, including the time required to collect, deliver, and

preprocess the CML measurements.

Beyond the proposed operational validation, one direction is to move from single-time reconstruction to longer event sequences over the same region and fixed link network. Even without changing the objective, this would show whether the method behaves stably over time or produces frame-to-frame artifacts when applied independently to consecutive observations. A second step would be to incorporate temporal structure directly. In meteorology, *optical flow* refers to methods that estimate an apparent motion field between consecutive images (Ayzel et al., 2019). For CML-based nowcasting, each set of link measurements at a discrete time would first be reconstructed into a rainfall map, and a sequence of these reconstructed maps would be used to estimate the motion field. Advecting the latest reconstructed map according to this field would provide a short-term forecast. Once the next set of CML measurements becomes available, the advected map could serve as a temporal prior: the reconstruction would be encouraged to remain close to the forecast, while the attenuation-mismatch term would allow it to depart where the new observations indicate that the forecast is inaccurate.

A further possibility would be to adapt the present reconstruction framework to radar-CML data fusion, building on existing variational fusion work (Bianchi et al., 2013). Radar observations could enter through a radar-mismatch term, combining radar’s spatial coverage with the path-integrated, near-ground information provided by CMLs. This adaptation cannot be evaluated independently with the current benchmark because its EU-RADCLIM maps already serve as ground truth; an operational test would require real radar and CML measurements, verification data withheld from the solver, and careful calibration of the relative observation weights.

Another possible direction is to compare the deterministic solvers studied here with learning-based inverse approaches. Recent work, for example, has studied CML rain-field reconstruction using a Bayesian inverse formulation with learned generative priors (Moufadi et al., 2026). Such methods could provide a useful point of comparison, but they would require a separate evaluation of training data, prior assumptions, and generalization beyond the benchmark setting used here.

Whether the improvements over IDW observed under controlled, model-consistent conditions persist with operational CML data remains the central question for future validation.

Appendix A

Supplementary Analyses

A.1 Largest-Patch Link-Network Summary

This appendix summarizes the link-network properties of the largest evaluated patch shown in Figure 5. Selected by physical area, this is the 88 km by 88 km patch from January 29, 2023 at 19:00, with an area of 7744 km² and 1,894 retained links. The distance-bin map shows which pixels are close to or far from nearby links, but it does not by itself describe the link network that produced that geometry.

Here, link length refers to the path length of a link, carrier frequency is the microwave-link frequency in GHz, and endpoint degree is the number of links incident to a given endpoint node.

Link length matters because longer links average rainfall over longer paths and can smooth local rainfall variation. Carrier frequency matters because rain-induced attenuation depends on frequency. Endpoint degree gives a compact view of whether the network is mostly chain-like or whether many links meet at shared nodes.




Figure A.1: Histogram summaries for the largest patch. The left panel shows link path lengths, the middle panel shows carrier frequencies, and the right panel shows endpoint degrees, where endpoint degree is the number of links incident to a given endpoint node. These distributions provide context for the link geometry underlying the distance-bin map in the main text.

Of the 1,894 links, 73.4% are shorter than 4 km and 93.9% are shorter than 8 km; 69.0% operate between 35 and 40 GHz. Among the 1,123 distinct endpoints, 86.7% have degree below five. Thus, the network is dominated by short, 35–40 GHz links, and most endpoints have low degree.

A.2 Case Studies

Across the patches inspected manually, two recurring qualitative patterns emerged. In the first, Solver(GT) remains close to the ground-truth rainfall field while Convex Solver differs visibly. In the second, Solver(GT) and Convex Solver produce similar reconstructions, both visibly different from the ground-truth rainfall field. The two cases below illustrate these patterns.

The first example illustrates the first pattern. In the Black Sea coast, Romania case study, Solver(GT) closely resembles the ground-truth rainfall field, while Convex Solver gives a visibly different reconstruction. Table A.1 compares their final fields using J_{atten} and RMSE.

Solver	J_{atten}	RMSE (mm/h)
Solver(GT)	1.81×10^{-5}	0.102
Convex Solver	4.06×10^{-6}	0.702

Table A.1: Shared post-hoc evaluation measures for the final Solver(GT) and Convex Solver fields in the Black Sea coast, Romania case study.

For a complementary within-solver view, Table A.2 reports Solver(GT)’s weighted objective at initialization and after reconstruction.

Solver(GT) state	Solver(GT) weighted objective	RMSE (mm/h)
Ground-truth initialization	2.832	0.000
Final reconstruction	1.485	0.102

Table A.2: Solver(GT) weighted objective and RMSE at initialization and after reconstruction in the Black Sea coast, Romania case study.

Convex Solver has the lower J_{atten} , but Solver(GT) is substantially closer to the radar-derived ground-truth field. Thus, stronger consistency with the link observations does not necessarily imply lower pixelwise reconstruction error. Moreover, Solver(GT)’s increasing RMSE as its own objective decreases shows that this distinction also arises within a single formulation.




Figure A.2: Black Sea coast, Romania case study observed on January 28, 2023 at 09:00. The panel shows the ground truth, IDW, Solver(GT), and Convex Solver.




Figure A.3: Solver-specific objective traces for the Black Sea coast, Romania case study. Values can be compared over iterations within each panel, but not directly between panels because the solvers use different link models and objective normalizations.

The traces show that most of the objective reduction occurs early. For Convex Solver, the objective continues to decrease after several hundred iterations, but the visible change after the first few hundred iterations is small and is driven mainly by the first-difference penalty J_{1d} . Solver(GT) stops after 47 iterations when the ftol criterion is met, following the initial objective reduction.

The second example illustrates the second pattern. In the Northern Norway case study, Solver(GT) and Convex Solver closely resemble each other and have similar J_{atten} and RMSE values, as shown in Table A.3.

Solver	J_{atten}	RMSE (mm/h)
Solver(GT)	1.24×10^{-6}	0.552
Convex Solver	2.50×10^{-6}	0.565

Table A.3: Shared post-hoc evaluation measures for the final Solver(GT) and Convex Solver fields in the Northern Norway case study.

Solver(GT)’s own trajectory is summarized in Table A.4.

Solver(GT) state	Solver(GT) weighted objective	RMSE (mm/h)
Ground-truth initialization	3.142	0.000
Final reconstruction	0.289	0.552

Table A.4: Solver(GT) weighted objective and RMSE at initialization and after reconstruction in the Northern Norway case study.

Here Solver(GT) is slightly better than Convex Solver on both common evaluation measures, although the differences are small. As in the first case, however, Solver(GT) moves away from its exact ground-truth initialization while lowering its own weighted objective.




Figure A.4: Northern Norway case study observed on January 31, 2023 at 09:00. The panel shows the ground truth, IDW, Solver(GT), and Convex Solver.

In this case, the error heatmaps show smaller areas of positive error for Solver(GT) and Convex Solver than for IDW.




Figure A.5: Solver-specific objective traces for the Northern Norway case study. As above, objective values should be interpreted within each solver rather than compared directly across panels.

As in the Black Sea case, the first few hundred Convex Solver iterations account for most of the visible objective reduction, with the remaining gradual decrease driven mainly by the first-difference penalty J_{1d} .

Taken together, these appendix case studies support the main-text interpretation that agreement with the link observations and pixelwise agreement with the radar-derived ground truth capture different aspects of reconstruction quality. In the first case, Convex Solver has lower J_{atten} but substantially higher RMSE than Solver(GT); in the second, the two methods are close on both measures and produce similar reconstructions. In both cases, Solver(GT) reduces its own objective while increasing its pixelwise error from zero.

Statement on the Usage of Generative Digital Assistants

For this thesis, I used OpenAI ChatGPT and OpenAI Codex as generative digital assistants. Their use included support for programming and debugging Python and LaTeX code, rendering and checking report artifacts, converting numerical summaries into LaTeX tables, comparing generated figures and output files, and revising explanatory prose for clarity, structure, and readability.

The tools were used as assistants rather than as independent scientific sources. Mathematical formulations, empirical results, numerical values, figures, citations, and final interpretations were checked against the project code, generated artifacts, local files, and the cited literature. I reviewed and edited the thesis text and take full responsibility for the submitted content, including the accuracy and completeness of its citations and references.

I am aware of the potential risks of using generative digital assistants, including incorrect suggestions, hallucinated claims, and over-polished wording that may obscure uncertainty. For this reason, I used them with caution and critical review.

References

- Eugene L. Allgower and Kurt Georg. *Introduction to Numerical Continuation Methods*. SIAM, Philadelphia, 2003. doi: 10.1137/1.9780898719154.
- David Atlas and Carlton W. Ulbrich. Path- and area-integrated rainfall measurement by microwave attenuation in the 1–3 cm band. *Journal of Applied Meteorology*, 16(12):1322–1331, 1977. doi: 10.1175/1520-0450(1977)016<1322:PAAIRM>2.0.CO;2.
- Georgy Ayzel, Maik Heistermann, and Tanja Winterrath. Optical flow models as an open benchmark for radar-based precipitation nowcasting (rainymotion v0.1). *Geoscientific Model Development*, 12(4):1387–1402, 2019. doi: 10.5194/gmd-12-1387-2019.
- Blandine Bianchi, Peter Jan van Leeuwen, Robin J. Hogan, and Alexis Berne. A variational approach to retrieve rain rate by combining information from rain gauges, radars, and microwave links. *Journal of Hydrometeorology*, 14(6):1897–1909, 2013. doi: 10.1175/JHM-D-12-094.1.
- Richard H. Byrd, Peihuang Lu, Jorge Nocedal, and Ciyou Zhu. A limited memory algorithm for bound constrained optimization. *SIAM Journal on Scientific Computing*, 16(5):1190–1208, 1995. doi: 10.1137/0916069.
- Christian Chwala and Harald Kunstmann. Commercial microwave link networks for rainfall observation: assessment of the current status and future challenges. *WIREs Water*, 6(2):e1337, 2019. doi: 10.1002/wat2.1337.
- Adam Eshel, Jonatan Ostrometzky, Shani Gat, Pinhas Alpert, and Hagit Messer. Spatial reconstruction of rain fields from wireless telecommunication networks—scenario-dependent analysis of IDW-based algorithms. *IEEE Geoscience and Remote Sensing Letters*, 17(5):770–774, 2020. doi: 10.1109/LGRS.2019.2935348.
- Adam Eshel, Hagit Messer, Harald Kunstmann, Pinhas Alpert, and Christian Chwala. Quantitative analysis of the performance of spatial interpolation methods for rainfall estimation using commercial microwave links. *Journal of Hydrometeorology*, 22(4):831–843, 2021. doi: 10.1175/JHM-D-20-0164.1.
- D. Giuli, A. Toccafondi, G. Biffi Gentili, and A. Freni. Tomographic reconstruction of rainfall fields through microwave attenuation measurements. *Journal of Applied Meteorology*, 30(9):1323–1340, 1991. doi: 10.1175/1520-0450(1991)030<1323:TROFT>2.0.CO;2.
- Oren Goldshtein, Hagit Messer, and Artem Zinevich. Rain rate estimation using measurements from commercial telecommunications links. *IEEE Transactions on Signal Processing*, 57(4):1616–1625, 2009. doi: 10.1109/TSP.2009.2012554.
- Maximilian Graf, Christian Chwala, Julius Polz, and Harald Kunstmann. Rainfall estimation from a German-wide commercial microwave link network: Optimized processing and validation for 1 year of data. *Hydrology and Earth System Sciences*, 24(6):2931–2950, 2020. doi: 10.5194/hess-24-2931-2020.
- A. R. Holt, J. W. F. Goddard, G. J. G. Upton, M. J. Willis, A. R. Rahimi, P. D. Baxter, and C. G. Collier. Measurement of rainfall by dual-wavelength microwave attenuation. *Electronics Letters*, 36(25):2099–2101, 2000. doi: 10.1049/el:20001468.

- International Telecommunication Union. Recommendation ITU-R P.838-3: Specific attenuation model for rain for use in prediction methods. Technical Report ITU-R P.838-3, International Telecommunication Union, 2005. URL <https://www.itu.int/rec/R-REC-P.838-3-200503-I/en>.
- Hidde Leijnse, Remko Uijlenhoet, and J. N. M. Stricker. Rainfall measurement using radio links from cellular communication networks. *Water Resources Research*, 43(3):W03201, 2007. doi: 10.1029/2006WR005631.
- Bin Lian, Zhongcheng Wei, Xiang Sun, Zhihua Li, and Jijun Zhao. A review on rainfall measurement based on commercial microwave links in wireless cellular networks. *Sensors*, 22(12):4395, 2022. doi: 10.3390/s22124395.
- Jack L. McKee and Andrew D. Binns. A review of gauge-radar merging methods for quantitative precipitation estimation in hydrology. *Canadian Water Resources Journal*, 41(1-2):186–203, 2016. doi: 10.1080/07011784.2015.1064786.
- Hagit Messer, Artem Zinevich, and Pinhas Alpert. Environmental monitoring by wireless communication networks. *Science*, 312(5774):713, 2006. doi: 10.1126/science.1120034.
- Badr Moufad, Albina Ilina, Hai Victor Habi, Salem Lahlou, Yazid Janati, Hagit Messer, and Eric Moulines. Bayesian rain field reconstruction using commercial microwave links and diffusion model priors, 2026. URL <https://arxiv.org/abs/2605.05520>. arXiv:2605.05520.
- R. L. Olsen, D. V. Rogers, and D. B. Hodge. The aR^b relation in the calculation of rain attenuation. *IEEE Transactions on Antennas and Propagation*, 26(2):318–329, 1978. doi: 10.1109/TAP.1978.1141845.
- Jonas Olsson, Laura Horváth-Varga, Remco van de Beek, Maximilian Graf, Aart Overeem, Magdalena Szaton, Vojtěch Bareš, Nejc Bezak, Christian Chwala, Carlo De Michele, Martin Fencel, Jochen Seidel, and Andrijana Todorović. How close are opportunistic rainfall observations to providing societal benefit? *Journal of Hydrometeorology*, 26(11):1585–1602, 2025. doi: 10.1175/JHM-D-25-0043.1.
- Aart Overeem, Hidde Leijnse, and Remko Uijlenhoet. Country-wide rainfall maps from cellular communication networks. *Proceedings of the National Academy of Sciences*, 110(8):2741–2745, 2013. doi: 10.1073/pnas.1217961110.
- Aart Overeem, Hidde Leijnse, and Remko Uijlenhoet. Two and a half years of country-wide rainfall maps using radio links from commercial cellular telecommunication networks. *Water Resources Research*, 52(10):8039–8065, 2016. doi: 10.1002/2016WR019412.
- Aart Overeem, Bas Walraven, Hidde Leijnse, and Remko Uijlenhoet. Four-year commercial microwave link dataset for the netherlands. 4TU.ResearchData, Version 1, 2024. doi: 10.4121/be252844-b672-471e-8d69-27269a862ec1.v1.
- Aart Overeem, Hidde Leijnse, Else van den Besselaar, Gerard van der Schrier, and Jan Fokke Meirink. EURADCLIM: The european climatological gauge-adjusted radar precipitation dataset (1-h accumulations). KNMI Radar Team, Version 3.0, 2025. doi: 10.21944/1rxx-ev62.
- Julius Polz, Christian Chwala, Maximilian Graf, and Harald Kunstmann. Rain event detection in commercial microwave link attenuation data using convolutional neural networks. *Atmospheric Measurement Techniques*, 13(7):3835–3853, 2020. doi: 10.5194/amt-13-3835-2020.

- A. R. Rahimi, A. R. Holt, G. J. G. Upton, and R. J. Cummings. The use of dual-frequency microwave links for measuring path-averaged rainfall. *Journal of Geophysical Research: Atmospheres*, 108(D15):4467, 2003. doi: 10.1029/2002JD003202.
- Donald Shepard. A two-dimensional interpolation function for irregularly-spaced data. In *Proceedings of the 1968 23rd ACM National Conference*, pages 517–524. Association for Computing Machinery, 1968. doi: 10.1145/800186.810616.
- G. J. G. Upton, A. R. Holt, R. J. Cummings, A. R. Rahimi, and J. W. F. Goddard. Microwave links: The future for urban rainfall measurement? *Atmospheric Research*, 77(1–4):300–312, 2005. doi: 10.1016/j.atmosres.2004.10.009.
- Ciyu Zhu, Richard H. Byrd, Peihuang Lu, and Jorge Nocedal. Algorithm 778: L-BFGS-B: Fortran subroutines for large-scale bound-constrained optimization. *ACM Transactions on Mathematical Software*, 23(4):550–560, 1997. doi: 10.1145/279232.279236.
- Artem Zinevich, Pinhas Alpert, and Hagit Messer. Estimation of rainfall fields using commercial microwave communication networks of variable density. *Advances in Water Resources*, 31(11):1470–1480, 2008. doi: 10.1016/j.advwatres.2008.03.003.